

# Cloud-based software

Antonio Brogi

Department of Computer Science  
University of Pisa

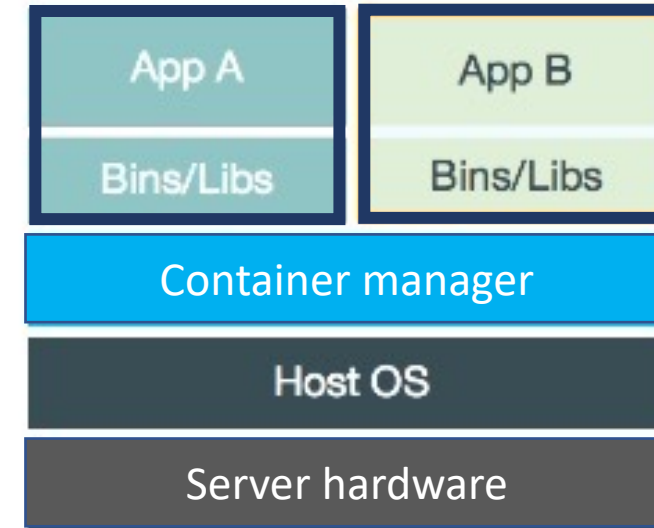
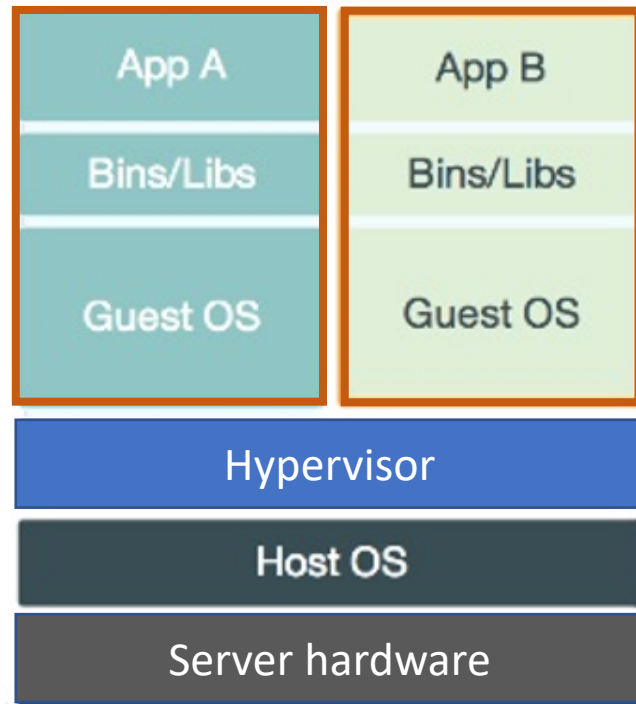


# Cloud computing

- Powerful computer hardware + high-speed networking have led to the development of cloud computing
- Virtualised resources (e.g. compute, storage, platform) accessible on demand
- Main advantages
  - Scalability (maintain performance as load increases)
  - Elasticity (adapt server configuration to changing demands)
  - Resilience (maintain service in the event of server failures)
  - Cost (CapEx reduction, pay-per-use)

# Virtualization and containers

# From VMs to containers



- Server virtualization
- Many virtual machines (each with its OS) running on same physical server
- ("Type 1" hypervisor loaded directly on HW)


Key idea of containers: Exploit OS kernel's capability of allowing multiple isolated user-space instances

- + lighter (require less resources, from Gb to Mb)
- + faster to start (from mins to secs)
- share same OS


# What is Docker?

Docker is a platform that allows us  
to run applications in an isolated environment

Docker allows us to develop and run portable applications  
by exploiting *containers*



# How old are containers?



- For decades, UNIX *chroot* command provided a simple form of filesystem isolation
- 1998 – FreeBSD *jail* utility extended *chroot* sandboxing to processes
- 2005 – Google started developing *CGroups* for Linux kernel and began moving its infrastructure to containers
- 2008 – Linux Containers (LXC) provided a complete containerization solution

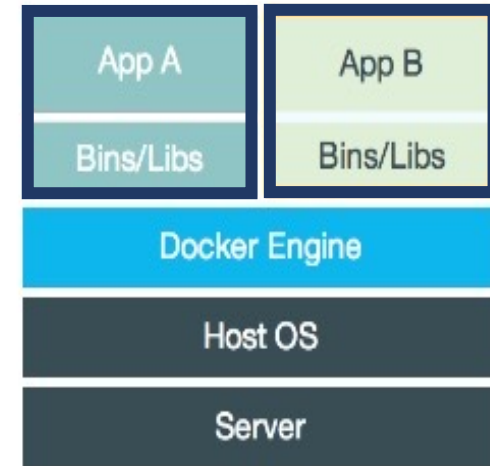
2013 - Docker added the missing pieces - *portable images* and *friendly UI* – to the containerization puzzle, and containers entered the mainstream

The Docker platform consisted of:

- **Docker Engine** (for creating and running containers)
- **Docker Hub** (for distributing containers)

# Docker

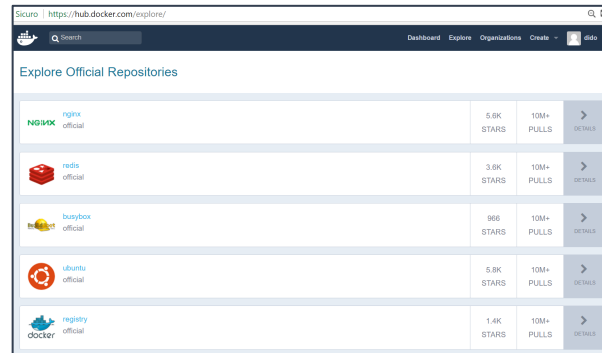
- Docker exploits container-based virtualization to run multiple isolated guest instances on the (same) OS
- Software components are packaged into ***images***, which are exploited as read-only templates to create and run ***containers***
- External ***volumes*** can be mounted to ensure data persistence




“Build, ship, and run any app, anywhere”

# Docker images

- (read-only) templates used to create containers
- stored in a (private or public) Docker **registry**
  - registry structured in repositories
  - each repository contains a set of images, for different versions of a software
  - images identified by pairs **repository:tag**
  - [Docker Hub](#)



- images structured into (read-only) layers
- each layer is in turn an image, lowest layer called base image
- running container can write in new top layer, changes can be committed into new image





# Docker commands



```
FROM node
LABEL maintainer ian.miehl@gmail.com
RUN git clone https://github.com/docker-in-practice/todo.git
WORKDIR todo
RUN npm install
RUN chmod -R 777 /todo
EXPOSE 8000
CMD ["npm","start"]
```

Dockerfile

build



image

pull


push



registry


commit

run



container

# Learn Docker in 12 min



<https://www.youtube.com/watch?v=YFI2mCHdv24>



## Advantages of Docker

- same environment
- sandbox projects
- it just works

## Containers

- lighter and faster to start than VMs
- share OS
- images defined via Dockerfile
- build and run

## Demo

install Docker

simple php “hello world” app

index.php

```
1 <?php
2
3 echo "Hello, World"
```

Dockerfile

find php image from Docker Hub

```
1 FROM php:7.0-apache
2 COPY src/ /var/www/html
3 EXPOSE 80
```

>docker build -t hello-world .

>docker run -p 80:80 hello-world

using a volume to share folder between host and container  
one process per container

# Docker Compose in 12 min



<https://www.youtube.com/watch?v=Qw9zIE3t8Ko>



//product service

api.py

```

3 from flask import Flask
4 from flask_restful import Resource, Api
5
6 app = Flask(__name__)
7 api = Api(app)
8
9 class Product(Resource):
10     def get(self):
11         return {
12             'products': ['Ice cream',
13                          'Chocolate',
14                          'Fruit',
15                          'Eggs']
16         }
17
18 api.add_resource(Product, '/')
19
20 if __name__ == '__main__':
21     app.run(host='0.0.0.0', port=80, debug=True)
  
```

requirements.txt

```

1 Flask==0.12
2 flask-restful==0.3.5
  
```

Dockerfile

```

1 FROM python:3-onbuild
2 COPY . /usr/src/app
3 CMD ["python", "api.py"]
  
```

docker-compose.yml

```

1 version: '3'
2
3 services:
4   product-service:
5     build: ./product
6     volumes:
7       - ./product:/usr/src/app
8     ports:
9       - 5001:80
10
11   website:
12     image: php:apache
13     volumes:
14       - ./website:/var/www/html
15     ports:
16       - 5000:80
17     depends_on:
18       - product-service
  
```

>docker-compose up  
>docker ps  
>docker-compose stop

//website

index.php

```

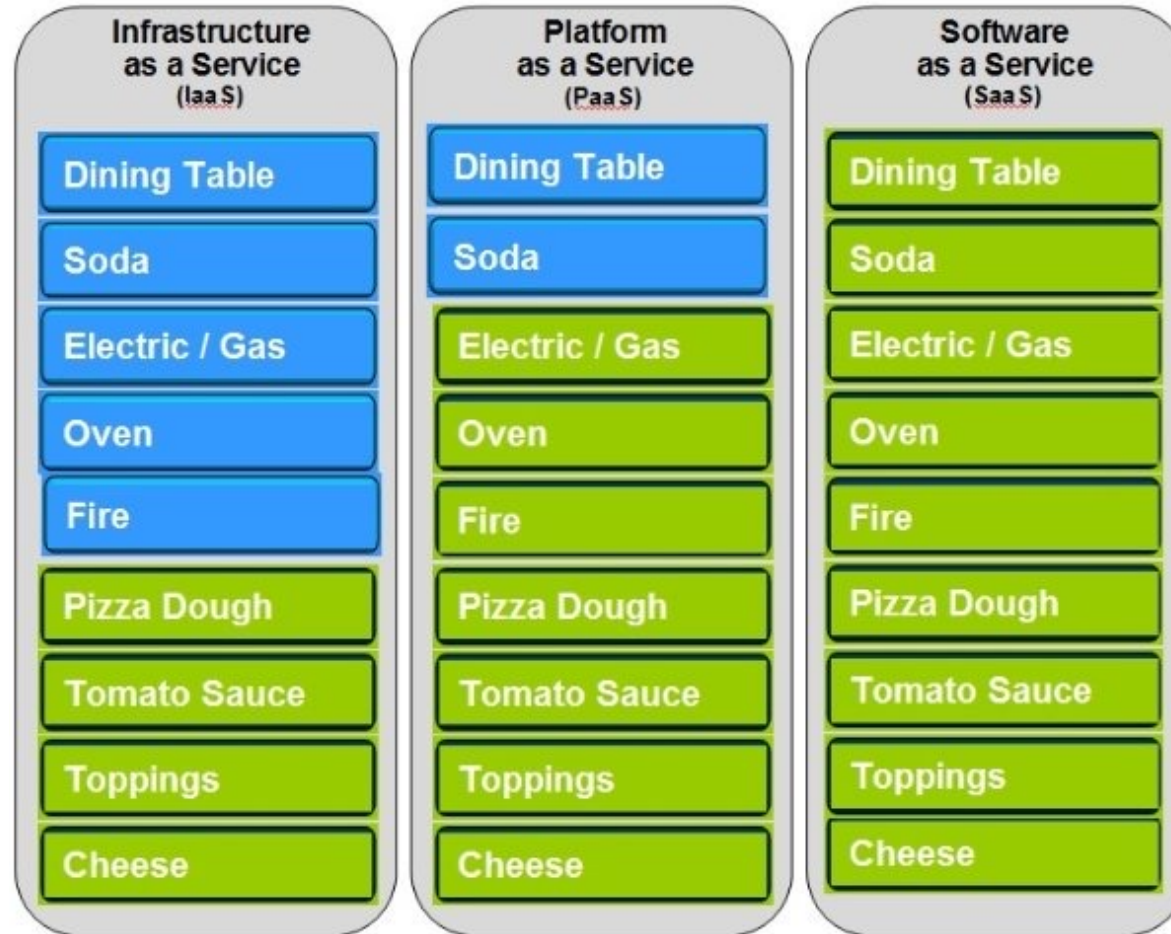
1 <html>
2   <head>
3     <title>My Shop</title>
4   </head>
5
6   <body>
7     <h1>Welcome to my shop</h1>
8     <ul>
9       <?php
10         $json = file_get_contents('http://product-service');
11         $obj = json_decode($json);
12
13         $products = $obj->products;
14         foreach ($products as $product) {
15           echo "<li>$product</li>";
16         }
17       ?>
18     </ul>
19   </body>
20 </html>
  
```



Virtualization and containers  
Everything as a service



# Pizza as a Service



- You Manage
- Vendor Manages



# \*aaS

+ FaaS ...

- SaaS provides software on-demand for use, accessible via thin clients or APIs
- SaaS provider manages infrastructure + OS + app
- Client responsible for nothing
- Example: salesforce.com

- PaaS provides whole platform as a service (VMs, OS, services, SDKs,...)
- PaaS provider manages infrastructure + OS + enabling SW
- Client responsible for installing and managing app
- Examples: Heroku, Azure, GAE

- IaaS provides (virtualized) servers, storage, networking
- IaaS provider manages all infrastructure
- Client responsible for all other aspects of the deployment (e.g., OS, app)
- Example: EC2, S3



Virtualization and containers  
Everything as a service  
SaaS

# SaaS

Software products were initially installed on customers' computers

- Customers had to configure software
- Customers had to deal with software updates
- Software product company had to maintain different product versions

## SaaS

- Product delivered as a service
- Customers do not have to install software
- They pay a subscription and access product from remote

# Benefits of SaaS for product providers

## Cash flow

Regular cash flow: Customers pay periodic subscription or pay-per-usage

## Update management

You control product updates, customers receive update at same time

No need of simultaneously maintaining several versions

Reduced costs

## Continuous deployment

You can deploy new software versions as soon as changes have been made and tested

## Payment flexibility

Different payment options can attract wider range of customers

e.g. small companies or individuals can avoid paying large upfront software costs

## Try before you buy

You can make early free/low-cost product available and get customer feedback

## Data collection

You can easily collect data on product usage and customers

# Pros & cons for customers

- **Mobile, laptop, and desktop access**
- **No upfront costs for software or servers**
- **Immediate software updates**
- **Reduced software management costs**
- **Privacy regulation conformance**
  - e.g. EU countries have strict laws on storage of personal info
- **Security concerns**
  - Customers may not want to pass data control to external provider
- **Network constraints**
  - Can limit response time when much data transfer
- **Data exchange**
  - Can be difficult if cloud does not provide suitable API
- **Loss of control over updates**
- **Service lock-in**

# Design issues for SaaS

## **Local vs. remote processing**

Some features can be executed locally

- + reduced network traffic
- + increased response speed
- increased power consumption for battery-powered devices

## **Authentication**

Product own authentication system vs.

Federated authentication vs.

Individual Google/LinkedIn/... credentials

## **Information leakage**

Security risks with multiple users from multiple organizations

## **Multi-tenant vs. multi-instance database management**

Single repository vs. separate copies of system and database

Virtualization and containers

Everything as a service

SaaS

Multi-tenant and multi-instance systems

# Multi-tenant systems

- Single database schema shared by all system's users
- Items in database tagged with tenant identifier to provide “logical isolation”

Stock management					
Tenant	Key	Item	Stock	Supplier	Ordered
T516	100	Widg 1	27	S13	2017/2/12
T632	100	Obj 1	5	S13	2017/1/11
T973	100	Thing 1	241	S13	2017/2/7
T516	110	Widg 2	14	S13	2017/2/2
T516	120	Widg 3	17	S13	2017/1/24
T973	100	Thing 2	132	S26	2017/2/12

Advantages	Disadvantages
<b>Resource utilization</b> The SaaS provider has control of all the resources used by the software and can optimize the software to make effective use of these resources.	<b>Inflexibility</b> Customers must all use the same database schema with limited scope for adapting this schema to individual needs. I explain possible database adaptations later in this section.
<b>Security</b> Multi-tenant databases have to be designed for security because the data for all customers are held in the same database. They are, therefore, likely to have fewer security vulnerabilities than standard database products. Security management is also simplified as there is only a single copy of the database software to be patched if a security vulnerability is discovered.	<b>Security</b> As data for all customers are maintained in the same database, there is a theoretical possibility that data will leak from one customer to another. In fact, there are very few instances of this happening. More seriously, perhaps, if there is a database security breach, then it affects all customers.
<b>Update management</b> It is easier to update a single instance of software rather than multiple instances. Updates are delivered to all customers at the same time so all use the latest version of the software.	<b>Complexity</b> Multi-tenant systems are usually more complex than multi-instance systems because of the need to manage many users. There is, therefore, an increased likelihood of bugs in the database software.



# Multi-tenant systems

Mid-size and large businesses

rarely want to use generic multi-tenant software

often prefer a customized version adapted to their own requirements

**Table 5.6** Possible customizations for SaaS

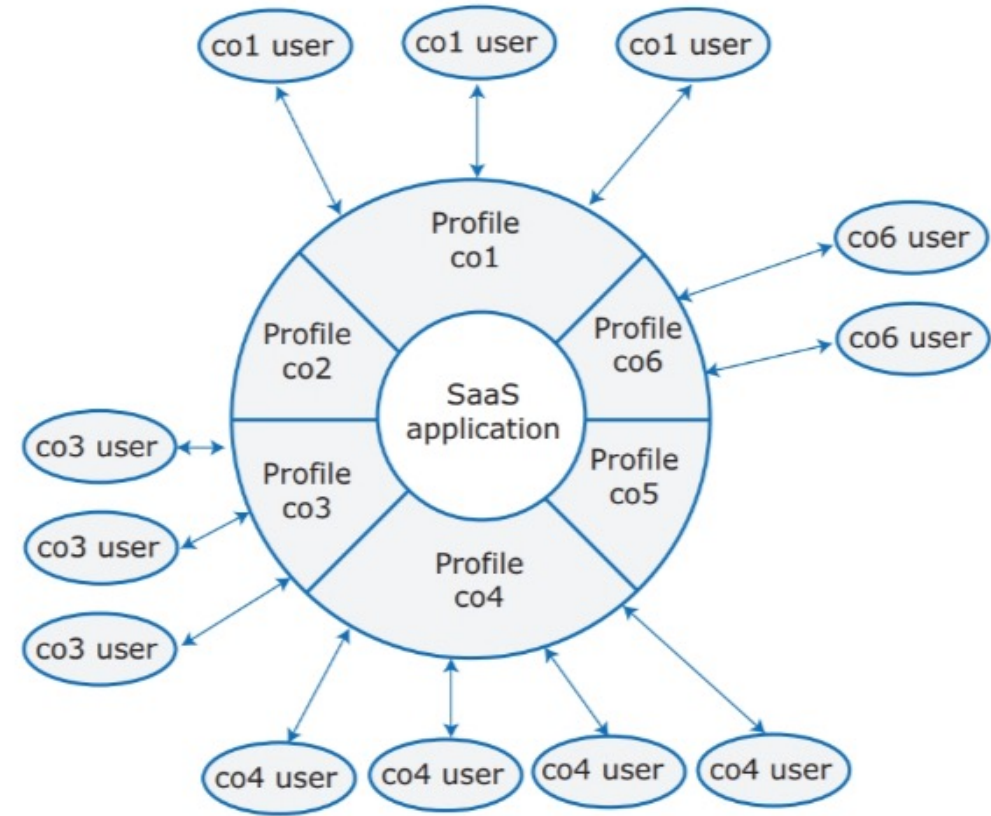
Customization	Business need
Authentication	Businesses may want users to authenticate using their business credentials rather than the account credentials set up by the software provider. I explain in Chapter 7 how federated authentication makes this possible.
Branding	Businesses may want a user interface that is branded to reflect their own organization.
Business rules	Businesses may want to be able to define their own business rules and workflows that apply to their own data.
Data schemas	Businesses may want to be able to extend the standard data model used in the system database to meet their own business needs.
Access control	Businesses may want to be able to define their own access control model that sets out the data that specific users or user groups can access and the allowed operations on that data.



# Multi-tenant systems: UI configurability

UI configurability can be implemented by employing user profiles

- Users asked to select their organization or provide their business email address
- Product uses profile information to create personalized version of interface
  - e.g. generic company name and logo replace with company's ones
  - e.g. some features of menus disabled



# Multi-tenant systems: DB schema adaptation

*Corporate users may wish to extend or adapt DB schema to meet their specific business needs*

***Solution I:*** Add a number of extra fields to each table and allow customers to use these fields as they wish

- Issues:
  - Difficult to know how many extra columns to include (too few will be insufficient, too many will lead to wasted space)
  - Different customers are likely to need different types of columns

Stock management								
Tenant	Key	Item	Stock	Supplier	Ordered	Ext 1	Ext 2	Ext 3
T516	100	Widg 1	27	S13	2017/2/12			
T632	100	Obj 1	5	S13	2017/1/11			
T973	100	Thing 1	241	S13	2017/2/7			
T516	110	Widg 2	14	S13	2017/2/2			
T516	120	Widg 3	17	S13	2017/1/24			
T973	100	Thing 2	132	S26	2017/2/12			

# Multi-tenant systems: DB schema adaptation

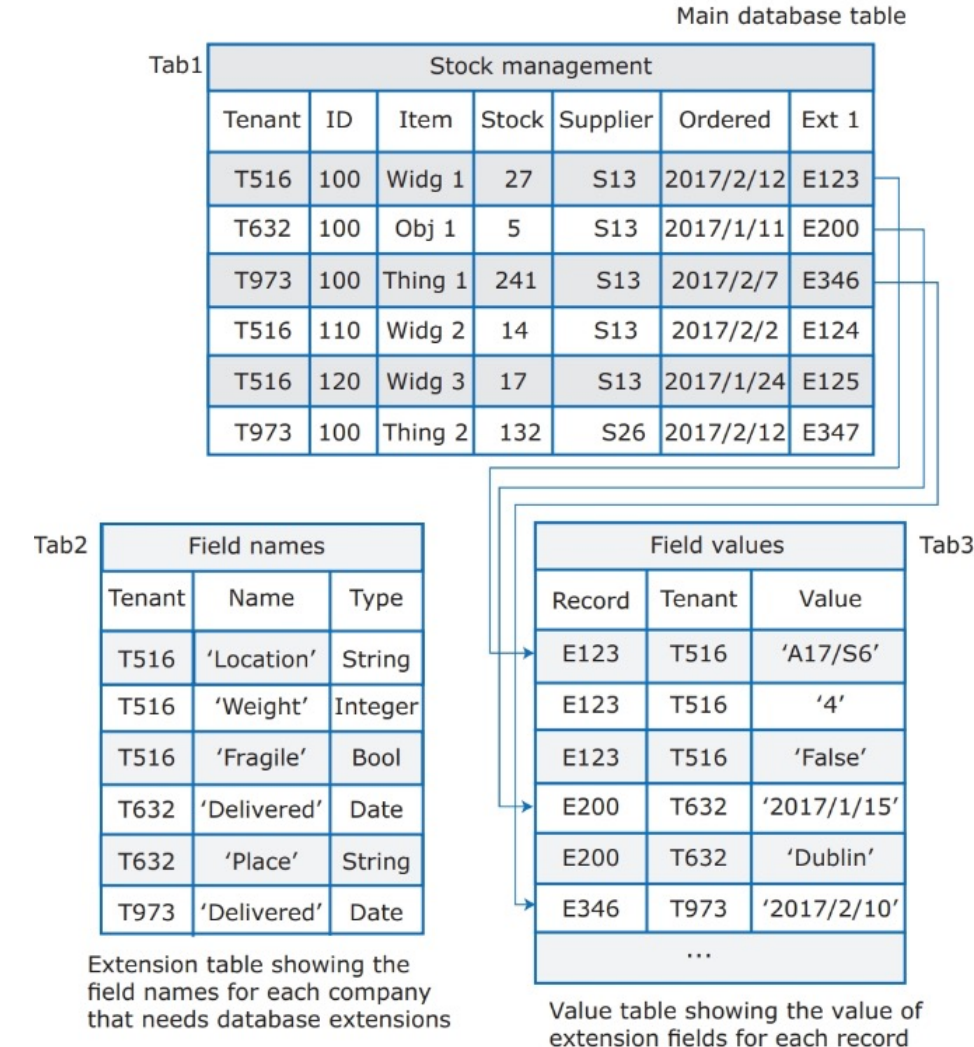
*Corporate users may wish to extend or adapt DB schema to meet their specific business needs*

**Solution II:** Add a field to each table that identifies a separate “extension table,” and allow customers to create these extension tables to reflect their needs.

e.g. The extension fields for tenant T516/ Item 100 are ‘Location’, ‘Weight’ and ‘Fragile’, with values ‘A17/S6’, ‘4’, and ‘False’

e.g. The extension fields for T632/Item 100 are ‘Delivered’ and ‘Place’, and their values are ‘2017/1/15’ and ‘Dublin’

Issue: added complexity



# Multi-tenant systems: Security issues

Security is the major concern of corporate customers with multi-tenant databases

- As information from all customers is stored in the same database, a software bug or an attack could lead to the data of some or all customers being exposed to others

## **Multilevel access control**

- access to data must be controlled at both the organizational level and the individual level
  - organizational level access control to ensure that any database operations act on only that organization's data
  - individual users accessing the data should also have their own access permissions

## **Encryption of data**

- To ensure corporate users that their data cannot be viewed by people from other companies if some kind of system failure occurs
- Encrypted data can be decrypted only when accessed with the appropriate key
- Encryption/decryption are computationally intensive operations → usually only sensitive data are encrypted

# Multi-instance systems

Each customer has its own system that is adapted to its needs, including its own database and security controls

- conceptually simpler than multi-tenant systems
- avoid security concerns such as inter-organization data leakage

# Multi-instance systems

Two types of multi-instance systems:

- ***VM-based***

- software instance and database for each customer run in its own VM
- all users from same customer may access the shared system database

- ***Container-based***

- each user has an isolated version of software and database running in a set of containers
- most suited to products in which users mostly work independently, with little data sharing (e.g. for individual users)

It is also possible to run containers on a virtual machine  
a business could have its own VM-based system and run containers on top of this for individual users



# Pros and cons of multi-instance databases

Advantages	Disadvantages
<p><b>Flexibility</b> Each instance of the software can be tailored and adapted to a customer's needs. Customers may use completely different database schemas and it is straightforward to transfer data from a customer database to the product database.</p>	<p><b>Cost</b> It is more expensive to use multi-instance systems because of the costs of renting many VMs in the cloud and the costs of managing multiple systems. Because of the slow startup time, VMs may have to be rented and kept running continuously, even if there is very little demand for the service.</p>
<p><b>Security</b> Each customer has its own database so there is no possibility of data leakage from one customer to another.</p>	<p><b>Update management</b> Many instances have to be updated so updates are more complex, especially if instances have been tailored to specific customer needs.</p>
<p><b>Scalability</b> Instances of the system can be scaled according to the needs of individual customers. For example, some customers may require more powerful servers than others.</p>	
<p><b>Resilience</b> If a software failure occurs, this will probably affect only a single customer. Other customers can continue working as normal.</p>	

Virtualization and containers

Everything as a service

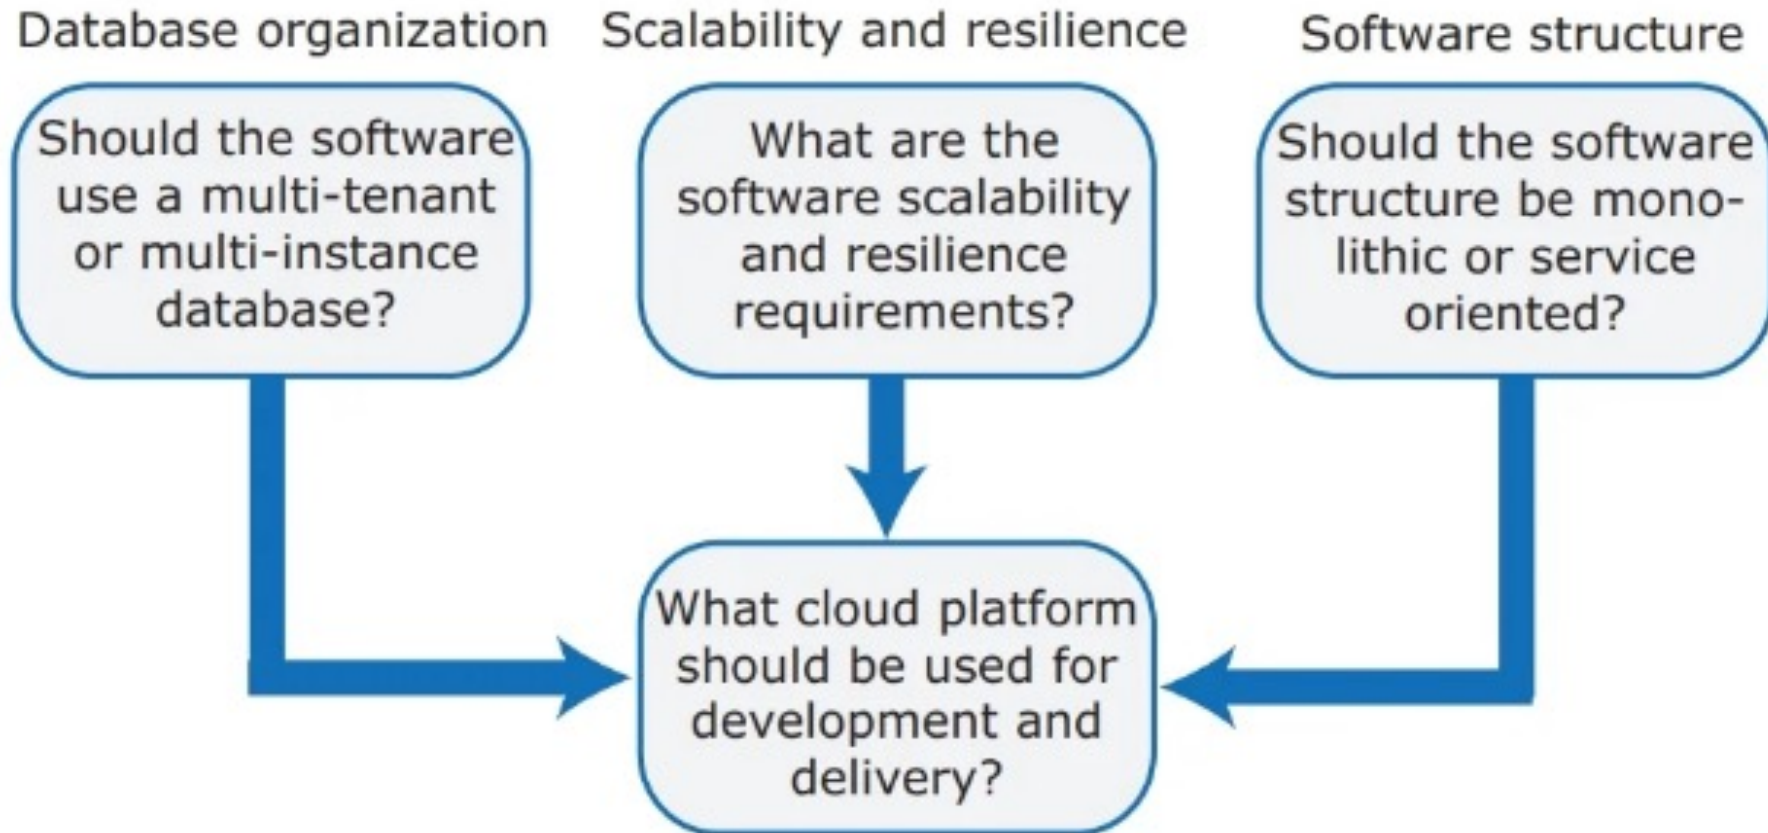
SaaS

Multi-tenant and multi-instance systems

Cloud software architecture



# Architectural decisions



# DB organization

There are three possible ways of providing a customer database in a cloud-based system:

1. As a multi-tenant system, shared by all customers for your product. This may be hosted in the cloud using large, powerful servers.
2. As a multi-instance system, with each customer database running on its own virtual machine.
3. As a multi-instance system, with each database running in its own container. The customer database may be distributed over several containers.

Factor	Key questions
Target customers	Do customers require different database schemas and database personalization? Do customers have security concerns about database sharing? If so, use a multi-instance database.
Transaction requirements	Is it critical that your products support ACID transactions where the data are guaranteed to be consistent at all times? If so, use a multi-tenant database or a VM-based multi-instance database.
Database size and connectivity	How large is the typical database used by customers? How many relationships are there between database items? A multi-tenant model is usually best for very large databases, as you can focus effort on optimizing performance.
Database interoperability	Will customers wish to transfer information from existing databases? What are the differences in schemas between these and a possible multi-tenant database? What software support will they expect to do the data transfer? If customers have many different schemas, a multi-instance database should be used.
System structure	Are you using a service-oriented architecture for your system? Can customer databases be split into a set of individual service databases? If so, use containerized, multi-instance databases.

# DB organization

Different types of customers have different expectations about software products

## Examples:

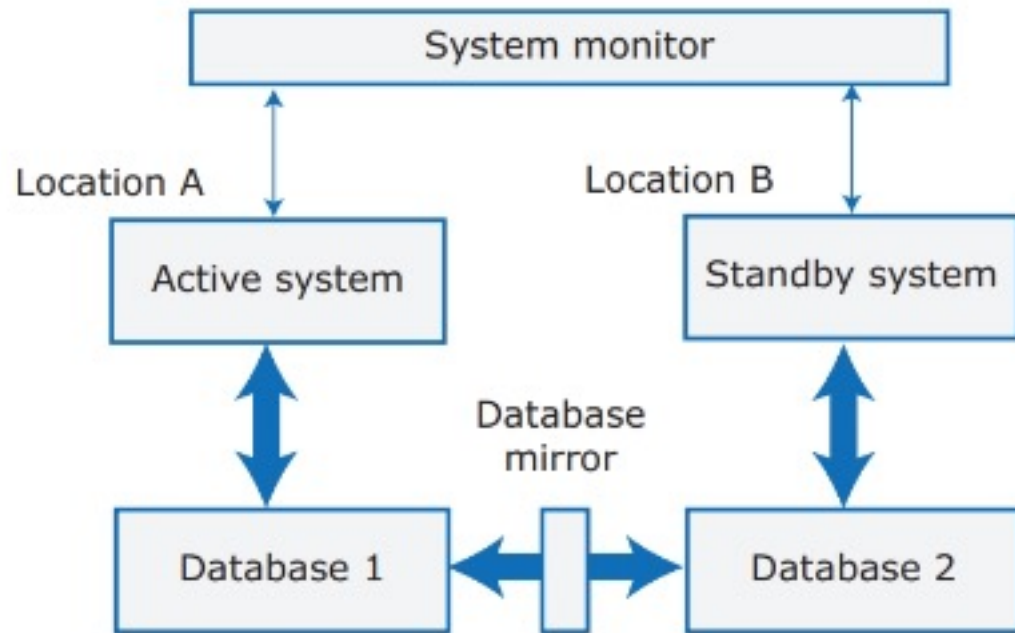
- Consumers or small businesses do not expect branding and personalization, local authentication system, or varying individual permissions → You can use a multi-tenant database with a single schema
- Large companies are more likely to want a database adapted to their needs → Possible to some extent with a multi-tenant system, easier with a multi-instance database
- For products in which database has to be consistent at all times (e.g. finance) → You need a transaction-based system: either a multi-tenant database or a database per customer running on a virtual machine (with all users from each customer sharing the VM-based database)
- ...

# Scalability

- Scalability = ability to adapt automatically to load changes
- Achieved by
  - By adding new virtual servers (scaling out)
  - Or by increasing the power of a system server (scaling up)in response to increasing load
- Scaling out typical of cloud-based systems
  - Product must be organized so that individual software components can be replicated and run in parallel
  - Load-balancing mechanisms direct requests to different instances of the components (e.g. with PaaS)

# Resilience

Resilience = ability to continue to deliver critical services in the event of system failure or malicious system use



## “Hot standby”

- Replicas of software and data maintained in different locations
- Database updates are mirrored
- System monitor continuously checks system status

## Cheaper alternative: “cool stanby”

- data restored from backup
- system unavailable till data restore is complete

# Software structure

Monolith vs. Fine-grain, stateless services

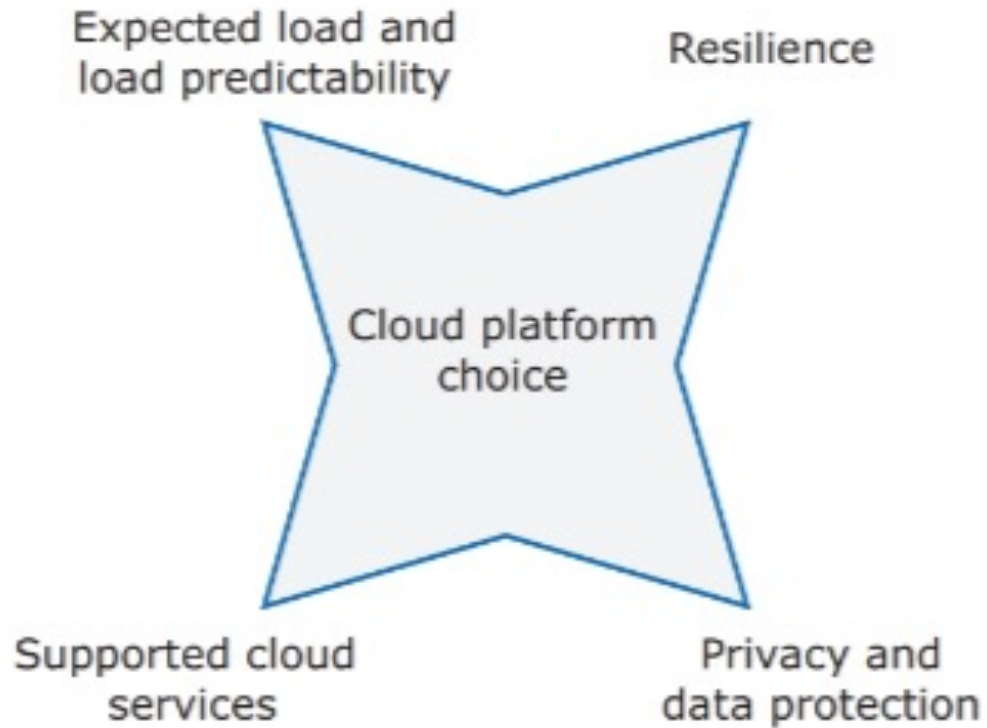
- + independent services that can be replicated, distributed, and migrated
- + particularly suitable for cloud-based software with services deployed in containers

Monolithic approach often employed to build prototype

When different parts have to be updated at different times and when only some parts have to scale → microservices

# Cloud platform

Technical issues in cloud platform choice



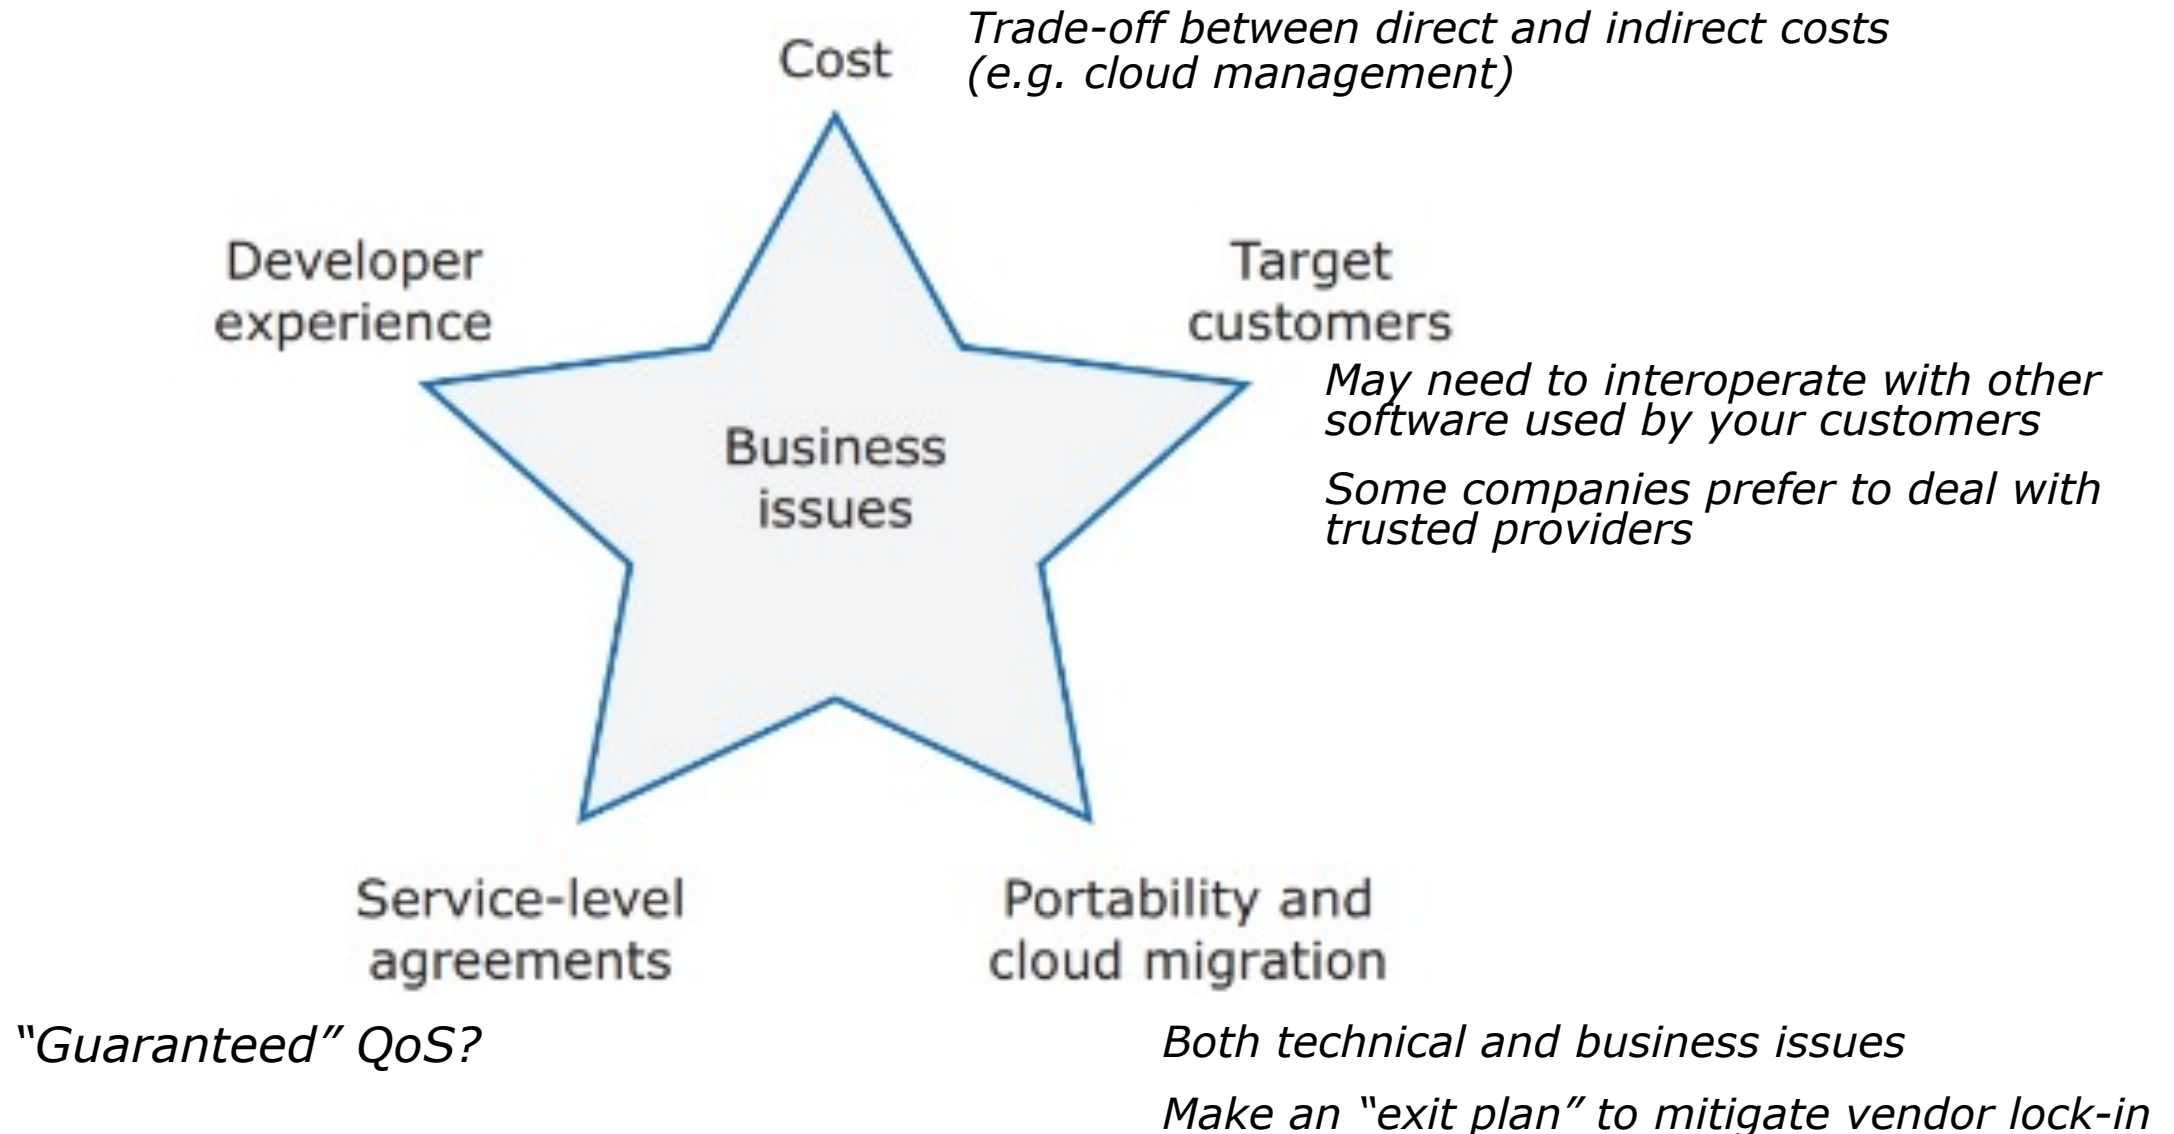
*Both technical and business issues*

*e.g. EU countries have strict requirements on data protection and on where data are stored → need cloud provider providing guarantees on storage locations*



# Cloud platform

## Business issues in cloud platform choice





# Reference



## Chapter 5 – Cloud-based software